

Parallel Debugging

Matthias Müller, Pavel Neytchev, Rainer Keller, Bettina Krammer

University of Stuttgart

High-Performance Computing-Center Stuttgart (HLRS)

www.hlrs.de

Outline

- Motivation
- Approaches and Tools
 - Memory Tracing Tools
 - **Valgrind**
 - MPI-Analysis Tools
 - **Marmot**
 - Debuggers
 - **TotalView**
 - **DDT**

Motivation - Problems of Parallel Programming

- All problems of serial programming
- Additional problems:
 - Increased difficulty to verify correctness of program
 - Increased difficulty to debug N parallel processes
 - New parallel problems:
 - **deadlocks**
 - **race conditions**
 - **irreproducibility**

Tools and Techniques to Avoid and Remove Bugs

- Programming techniques
- Static Code analysis
 - Compiler (with `-Wall` flag or similar), lint
- Runtime analysis
 - Memory tracing tools
 - Special OpenMP tools (assure, thread checker)
 - Special MPI tools
- Post mortem analysis
 - Debuggers

What is a Debugger?

- Common Misconception:
A debugger is a tool to find and remove bugs
- A debugger does:
 - tell you where the program crashed
 - help to gain a better understanding of the program and what is going on
- Consequence:
 - A debugger does not help much if your program does not crash, e.g. just gives wrong results
 - Avoid using a debugger as far as possible.
 - Use it as last resort.

Programming Techniques

- Think about a verbose execution mode of your program
- Use a careful/paranoid programming style
 - check invariants and pre-requisites
(`assert(m >= 0)`, `assert(v < c)`)

Static Code Analysis – Compiler Flags

- Use the debugging/assertion techniques of the compiler
 - use debug flags (-g), warnings (-Wall)
 - array bound checks in Fortran
 - use memory debug libraries (-lefence)

Avoiding Debuggers

- Write portable programs
 - it avoids future problems
 - **architectures/platforms have a short life**
 - **all compilers and libraries have bugs**
 - **all languages and standards include implementation defined behavior**
 - running on different platforms and architectures significantly increases the reliability
- Use verification tools for parallel programming like assure

Valgrind – Debugging Tool

Rainer Keller

University of Stuttgart

High-Performance Computing-Center Stuttgart (HLRS)

<http://www.hlrs.de>

Parallel Debugging

Müller, Neytchev, Keller, Krammer
Höchstleistungsrechenzentrum Stuttgart



Valgrind – Overview

- An Open-Source Debugging & Profiling tool.
- Works with any dynamically linked application.
- Emulates CPU, i.e. executes instructions on a synthetic x86.
- Currently it's only available for Linux/IA32.
- Prevents Error-swamping by suppression-files.
- Has been used on many **large** Projects:
KDE, Emacs, Gnome, Mozilla, OpenOffice.
- It's easily configurable to ease debugging & profiling through *skins*:
 - **Memcheck**: Complete Checking (every memory access)
 - Addrcheck: 2xFaster (no uninitialized memory check).
 - Cachegrind: A memory & cache profiler
 - **Callgrind**: A Cache & Call-tree profiler.
 - Helgrind: Find Races in multithreaded programs.

Valgrind – Usage

- Programs should be compiled with
 - Debugging support (to get position of bug in code)
 - Possibly without Optimization (for accuracy of position & less false positives):

```
gcc -O0 -g -o test test.c
```

- Run the application as normal as parameter to valgrind:

```
valgrind ./test
```

- Then start the MPI-Application as with TV as debugger:

```
mpirun -dbg=valgrind ./mpi_test
```

Valgrind – Memcheck

- Checks for:
 - Use of uninitialized memory
 - Malloc Errors:
 - **Usage of free'd memory**
 - **Double free**
 - **Reading/writing past malloced memory**
 - **Lost memory pointers**
 - **Mismatched malloc/new & free/delete**
 - Stack write errors
 - Overlapping arguments to system functions like `memcpy`.

Valgrind – Example 1/2

```
#define SIZE 10

int main (int argc, char * argv[])
{
    int size;
    int rank;
    int * array;
    MPI_Status status;
    array = malloc (SIZE * sizeof (int));

    if (rank == 0)
        MPI_Recv (array, SIZE+1, MPI_INT, 1, 4711, MPI_COMM_WORLD, &status, 0);
    else
        MPI_Send (array, SIZE+1, MPI_INT, 0, 4711, MPI_COMM_WORLD, 0);

    printf ("(Rank:%d) array[0]:%d\n", rank, array[0]);
    printf ("(Rank:%d) array[0]:%d\n", rank, array[0]);
    MPI_CHECK (MPI_Finalize ());
    return 0;
}
```

rusraink@pcgpc9:~/C/MPI_TESTS > █

Valgrind – Example 2/2

- PID • With Valgrind `mpirun -dbg=valgrind -np 2 ./mpi_murks:`

```
==11278== Invalid read of size 1
==11278==    at 0x4002321E: memcpy (../../memcheck/mac_replace_strmem.c:256)
==11278==    by 0x80690F6: MPID_SHMEM_Eagerb_send_short (mpich/../shmemshort.c:70)
.. 2 lines of calls to MPIch-functions deleted ...
==11278==    by 0x80492BA: MPI_Send (/usr/src/mpich/src/pt2pt/send.c:91)
==11278==    by 0x8048F28: main (mpi_murks.c:44)
==11278== Address 0x4158B0EF is 3 bytes after a block of size 40 alloc'd
==11278==    at 0x4002BBCE: malloc (../../coregrind/vg_replace_malloc.c:160)
==11278==    by 0x8048EB0: main (mpi_murks.c:39)
```

Buffer-Overrun by 4 Bytes in MPI_Send

```
....

==11278== Conditional jump or move depends on uninitialised value(s)
==11278==    at 0x402985C4: _IO_vfprintf_internal (in /lib/libc-2.3.2.so)
==11278==    by 0x402A15BD: _IO_printf (in /lib/libc-2.3.2.so)
==11278==    by 0x8048F44: main (mpi_murks.c:46)
```

Printing of uninitialized variable

- It can not find:
 - May be run with 1 process: One pending Recv (Marmot).
 - May be run with >2 processes: Unmatched Sends (Marmot).



Valgrind – Calltree 1/2

- The Calltree skin (like the cachegrind skin):
 - Tracks memory accesses to check Cache-hit/misses.
 - Additionally records call-tree information.

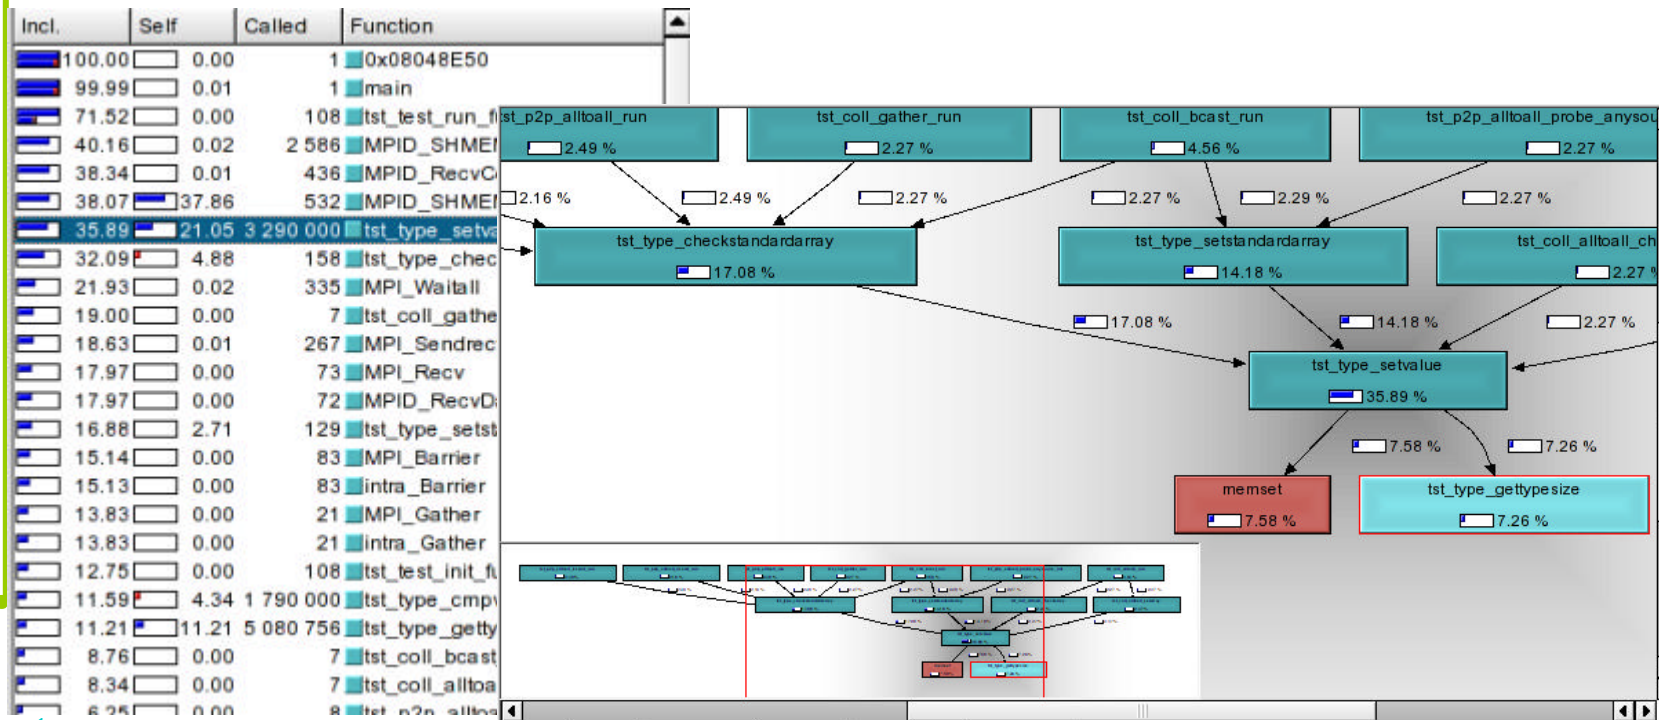
```
==11745== Calltree-0.9.6, a cache profiler for x86-linux.  
==11745== Copyright (C) 2002, and GNU GPL'd, by N.Nethercote and J.Weider  
==11745== Using valgrind-2.1.0, a program supervision framework for x86-]  
==11745== Copyright (C) 2000-2003, and GNU GPL'd, by Julian Seward.  
--11745-- warning: Pentium with 12 K micro-op instruction trace cache  
--11745--      Simulating a 16 KB cache with 32 B lines  
==11745== Estimated CPU clock rate is 1410 MHz  
==11745== For more details, rerun with: -v  
--11745--
```

- After the run, it reports overall program statistics:

```
==11810== D   refs:      497,790,574 (386,176,612 rd + 111,613,962 wr)  
==11810== D1  misses:      863,493 (   369,495 rd +   493,998 wr)  
==11810== L2d misses:    282,232 (   98,857 rd +   183,375 wr)  
==11810== D1  miss rate:      0.1% (      0.0% +      0.4% )  
==11810== L2d miss rate:      0.0% (      0.0% +      0.1% )
```

Valgrind – Calltree 2/2

- Even more interesting: the output trace-file.
- With the help of kcachegrind, one may:
 - Investigate, where Instr/L1/L2-cache misses happened2-2.
 - Which functions were called where & how often.



Valgrind – Deficiencies

- Valgrind cannot find of these Error-Classes:
 - Semantic Errors
 - Timing-critical errors
 - Uninitialised stack-memory not detected.
 - Problems with new instruction sets
(e.g. SSE/SSE2 is supported, certain Opcodes are **not**).
When using the Intel-Compiler: `-tpp5` for Pentium optimisation.

MARMOT

Bettina Krammer



University of Stuttgart

High-Performance Computing-Center Stuttgart (HLRS)

www.hlrs.de

Parallel Debugging

Müller, Neytchev, Keller, Krammer
Höchstleistungsrechenzentrum Stuttgart

H L R I S 

What is MARMOT?

- MPI analysis and checking tool to verify at runtime if an application conforms to the MPI standard.
- Library written in C++ that will be linked to the application.
- No source code modification of the application is required.
- Additional process working as debug server, i.e. the application will have to be run with `mpirun` for $n+1$ instead of n processes.
- Implementation of C and Fortran language binding of MPI-1.2 standard.
- Environment variables for tool behaviour and output (report of errors, warnings and/or remarks, trace-back, etc.).
- After the execution of the program the user can read a logfile to check for potential problems.

Availability of MARMOT

- Tests on different platforms, using different compilers and MPI implementations, e.g.
 - IA32/IA64 clusters (Intel, g++ compiler) mpich
 - IBM Regatta
 - NEC SX5
 - Hitachi SR8000
- Download and further information
<http://www.hlr.de/organization/tsc/projects/marmot/>

Example 1: request-reuse (source code)

```
/*
** Here we re-use a request we didn't free before
*/

#include <stdio.h>
#include <assert.h>
#include "mpi.h"

int main( int argc, char **argv ) {
    int size    = -1;
    int rank    = -1;
    int value   = -1;
    int value2  = -1;
    MPI_Status  send_status, recv_status;
    MPI_Request send_request, recv_request;

    printf( "We call Irecv and Isend with non-freed requests.\n" );
    MPI_Init( &argc, &argv );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    printf( " I am rank %d of %d PEs\n", rank, size );
```

Example 1: request-reuse (source code cont'd)

```
if( rank == 0 ){
    /** this is just to get the request used */
    MPI_Irecv( &value, 1, MPI_INT, 1, 18, MPI_COMM_WORLD, &recv_request );
    /** going to receive the message and reuse a non-freed request */
    MPI_Irecv( &value, 1, MPI_INT, 1, 17, MPI_COMM_WORLD, &recv_request );
    MPI_Wait( &recv_request, &recv_status );
    assert( value = 19 );
}
if( rank == 1 ){
    value2 = 19;
    /** this is just to use the request */
    MPI_Isend( &value, 1, MPI_INT, 0, 18, MPI_COMM_WORLD, &send_request );
    /** going to send the message */
    MPI_Isend( &value2, 1, MPI_INT, 0, 17, MPI_COMM_WORLD, &send_request );
    MPI_Wait( &send_request, &send_status );
}
MPI_Finalize();
return 0;
}
```

Example 1: request-reuse (output log)

We call Irecv and Isend with non-freed requests.

```
1 rank 0 performs MPI_Init
2 rank 1 performs MPI_Init
3 rank 0 performs MPI_Comm_size
4 rank 1 performs MPI_Comm_size
5 rank 0 performs MPI_Comm_rank
6 rank 1 performs MPI_Comm_rank
  I am rank 0 of 2 PEs
7 rank 0 performs MPI_Irecv
  I am rank 1 of 2 PEs
8 rank 1 performs MPI_Isend
9 rank 0 performs MPI_Irecv
10 rank 1 performs MPI_Isend
  ERROR: MPI_Irecv Request is still in use !!
11 rank 0 performs MPI_Wait
  ERROR: MPI_Isend Request is still in use !!
12 rank 1 performs MPI_Wait
13 rank 0 performs MPI_Finalize
14 rank 1 performs MPI_Finalize
```

Parallel Debuggers

Parallel Debuggers

- Most vendor debuggers have some support
- gdb has basic support for threads
- Debugging MPI programs with a “scalar” debugger is hard but possible
 - MPIch supports debugging with gdb attached to one process
 - manual attaching to the processes is possible
- TotalView is a good but expensive tool
- DDT is an alternative

TOTALVIEW



Parallel Debugging

Müller, Neytchev, Keller, Krammer
Höchstleistungsrechenzentrum Stuttgart



What is TotalView?

- Parallel debugger
- Source level debugging for C, C++, F77, F90, HPF
- **MPI, OpenMP**, Pthreads, PVM, shmem
- SMPs, MPPs, PVPs, Clusters
- Available on all major Unix Platforms and most Supercomputers
- GUI (independent of platform, exception Cray T3E)
 - TotalView 4.x based on tcl/tk
 - TotalView 5.x based on Motif

Availability of TotalView

- [Compaq Digital Alpha](#)
[HP-UX](#)
[IBM RS6000 and SP Power](#)
[SGI MIPS](#)
[Sun SPARC SunOS 5](#)
[Linux Intel IA32 \(RedHat\)](#)
[Linux Alpha \(RedHat\)](#)
Linux IA64
Linux Opteron
- Cray T3E by Cray
- Hitachi SR2201 by SofTek, SR8000
- NEC SX Series

Availability of TotalView at HWW

Platform	Availability	Remarks
Volvox	Yes	V6
Hitachi SR8000	Yes	V 4.0
Cray T3E	Yes	Cray 3.0.0
NEC SX	Yes	
SGI Onyx	No	Use cvd

Development
Platforms

More information:

<http://www.hlr.de/organization/tsc/services/tools/debugger/totalview>

Availability of TotalView at University Stuttgart

- Two user 8 CPU Floating License for University Stuttgart:

1. Download Software from <http://www.etnus.com>

2. Set environment variable for license:

`LM_LICENSE_FILE=7244@servint1.rus.uni-stuttgart.de`

More information about campus licenses available at

<http://www.hlrs.de/organization/tsc/services/tools/campus>

TotalView usage at HLRS

- Set USE_TOTALVIEW in your login scripts
- CRAY T3E: set USE_PROG_ENV
- Compile with -g compiler switch
CRAY T3E: compiler switch -G
- command name: **totalview**

Starting TotalView

On a new process:

```
% totalview myprog -a arguments to myprog
```

To debug MPI programs:

```
% totalview mpirun -a -nprocs 3 myprog
```

```
% mpirun -tv -np 3 myprog
```

To debug IBM POE programs:

```
% totalview poe -a myprog [args]
```

To debug CRAY T3E programs:

```
% totalview -X #procs myprog [args]
```


TotalView on Hitachi SR8000

- Compilation:
 - `f90 -g`
 - `cc -g`
 - `KCC -g --backend -tv`
- OpenMP
 - `f90 -g -omp -procnum=8`
 - `cc -g -omp -parallel=1 -O2`
- MPI
 - `mpirun -tv`

TotalView on HPN


- Compilation:
 - `f90 -g`
 - `cc -g`
 - `KCC -g`
- OpenMP
 - `gfortran -g`
 - `gcc -g`
 - `g++ -g`
- MPI
 - `mpirun -np #procs -tv ./a.out`

TotalView Exercise: Basic Look & Feel

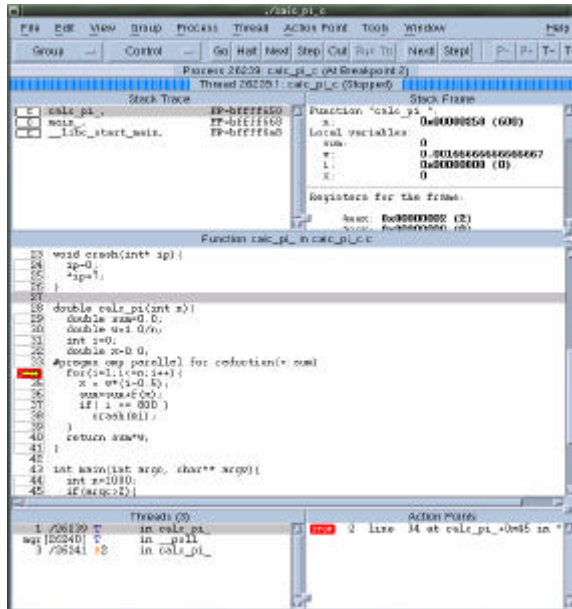
- Log into hwwhpn.hww.de
- Use bash as shell
- Change into directory
`~/TOTALVIEW/#NR/TOTALVIEW/SIMPLE`
- Compile `calc_pi_{f90,c,cc}.{f90,c,cc}`
- Start totalview with `totalview executable`

TotalView Windows



Root Window



Process Window




Data Windows

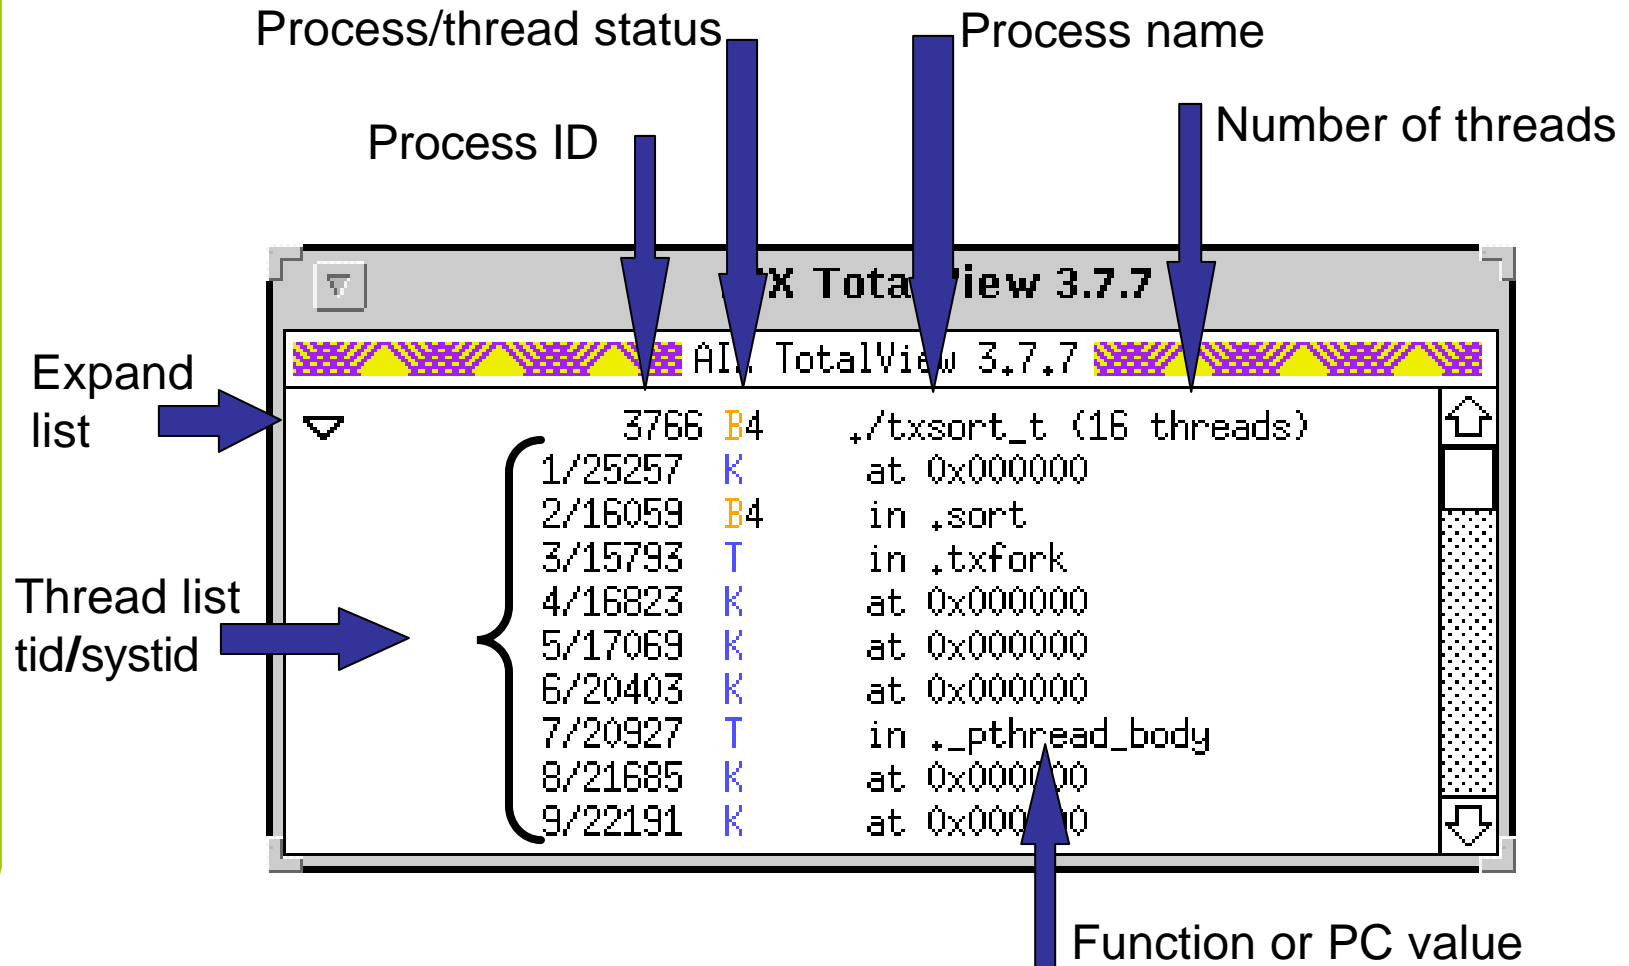


TotalView Mouse Buttons

- **Left** button is **Select**:
 - Chooses an item of interest,
 - Starts editing a item
- **Middle** button is **Dive**:
 - Gets more information about an item
 - **Shift+Dive** forces open a new window
- **Right** button is **Menu**:
 - Raises a menu of actions
 - All menus have a **Help** (^?) entry



TotalView Main Window



TotalView Process Window

Stack Trace
pane

Source
pane

Thread
pane

The screenshot shows the TotalView Process Window for the process `./calc_pi_c`. The window has a menu bar (File, Edit, View, Group, Process, Thread, Action Point, Tools, Window, Help) and a toolbar with buttons for Group, Control, Go, Halt, Next, Step, Out, Run To, Next!, Step!, and P-, P+, T-, T+.

The Stack Trace pane shows the following stack:

Frame	Address	FP
0	calc_pi_	FP=bffff650
1	main_	FP=bffff668
2	_libc_start_main	FP=bffff6a8

The Source pane shows the source code for the function `calc_pi` in `calc_pi.c`. The code is as follows:

```
23 void crash(int* ip){
24     ip=0;
25     *ip=7;
26 }
27
28 double calc_pi(int n){
29     double sum=0.0;
30     double w=1.0/n;
31     int i=0;
32     double x=0.0;
33     #pragma omp parallel for reduction(+:sum)
34     for(i=1;i<=n;i++){
35         x = w*(i-0.5);
36         sum=sum+f(x);
37         if( i == 800 )
38             crash(&i);
39     }
40     return sum*w;
41 }
42
43 int main(int argc, char** argv){
44     int n=1000;
45     if(argc>2){
```

The Thread pane shows the following threads:

Thread	Thread ID	Thread Name
1	/26239 T	in calc_pi_
mgr[26240]	T	in _poll
3	/26241 B2	in calc_pi_

The Action Points pane shows the following action points:

Action Point	Thread	Line	Code
STOP	2	line 34	at calc_pi_+0x45 in "

Process/thread
motion buttons


Local
variables
for the
selected
frame

Action
Points
pane


TotalView Source Pane

Gridded box is a possible site for a breakpoint

Select to set one



Current point of execution



Current function and source file



Function calc_pi_ in calc_pi_c.c

```
23 void crash(int* ip){
24     ip=0;
25     *ip=7;
26 }
27
28 double calc_pi(int n){
29     double sum=0.0;
30     double w=1.0/n;
31     int i=0;
32     double x=0.0;
33     #pragma omp parallel for reduction(+:sum)
34     for(i=1;i<=n;i++){
35         x = w*(i-0.5);
36         sum=sum+f(x);
37         if( i == 800 )
38             crash(&i);
39     }
```

- **Dive** on a source word to get more information
- **Select** a line to use **Run to selection** command
- **Select** or **dive** on a line number to set an action point



Parallel Debugging - Philosophy

- By default, TotalView places processes in groups
 - Program Group - Includes parent and all related processes
 - Share Group - Only processes that share the same source code
- Command can act on single process or share group
 - halt process (h) , halt group (H)
 - next step process (n), next step group (N)
 - go process (g), go group (G)

TotalView Exercise: Debug simple program

- Run calc_pi inside totalview:
 - Check where the program crashes
- Analyze core file with totalview
 - run calc_pi
 - execute `totalview calc_pi core`
- For advanced users: choose another programming paradigm:
 - MPI, OpenMP, MPI+OpenMP

TotalView support for debugging MPI

- Special support for MPI is available depending on your MPI library:
 - display message queue state of a process
- Supported MPI implementations:
 - mpich v1.1.0 or later (use -debug in configure)
 - HP MPI v1.6
 - Compaq MPI >v1.7
 - IBM, release >2.2 of Parallel Environment, threaded version of MPI
 - SGI MPI v1.3 or later

TotalView MPI Message Queue Window

Communicator name
and info

Non-blocking receive
operations

Unmatched incoming
messages

Non-blocking send
operations

- Dive on source or target
to refocus Process
window
- Dive on buffer to see
message contents

The screenshot shows the MPI Message Queue window for the process 'testsome.0'. The window title is 'testsome.0'. The main content area is titled 'Message State for "testsome.0" (1288,1)'. It displays the following information:

- MPI_COMM_WORLD**
 - Comm_size: 3
 - Comm_rank: 0
- Pending receives [0]**
 - Status: Pending
 - Source: 2 (testsome.2)
 - Tag: 0x00000000 (0)
 - User Buffer: 0x000605c0 -> 0x00000000 (0)
 - Buffer Length: 0x00000014 (20)
- Unexpected messages [0]**
 - Status: Complete
 - Source: 2 (testsome.2)
 - Tag: 0x00000002 (2)
 - System Buffer: 0x00000000
 - Buffer Length: 0x00000000 (0)
 - Received Length: 0x00000000 (0)
- Non-blocking sends [0]**
 - Status: Complete
 - Target: 2 (testsome.2)
 - Tag: 0x00000000 (0)
 - Buffer: 0x000605a0 -> 0x00000001 (1)
 - Buffer Length: 0x00000014 (20)
- MPI_COMM_WORLD_collective**
 - Comm_size: 3
 - Comm_rank: 0

TotalView Exercise: Parallel program

- Example in TOTALVIEW/MPI:
 - `deadlock_{c,cc,f90}.{c,cc,f90}`
 - start program with `mpirun -tv -np 2 a.out`
 - interrupt execution after “deadlock”
 - try to find the reason for the deadlock and fix it
- For advanced users:
 - `pending_{c,cc,f90}.{c,cc,f90}`
 - try to find pending message by setting breakpoint at `MPI_Finalize`

TotalView more information

- <http://www.etnus.com/products/totalview/index.html>
- <http://www.hlrs.de/organization/tsc/services/tools/debugger/totalview>
 - User Guide
 - Installation Guide
 - CLI Guide
 - Powerpoint Tutorial
- CRAY T3E: Online Documentation at <http://www.hlrs.de/platforms/crayt3e>

Distributed Debugging Tool (DDT)



Parallel Debugging

Müller, Neytchev, Keller, Krammer
Höchstleistungsrechenzentrum Stuttgart



What is DDT?

- Parallel debugger
- Source level debugging for C, C++, F77, F90
- MPI, OpenMP
- SMPs, Clusters
- Available on Linux distributions and Unix
- GUI (independent of platform, based on QT libraries)

Availability of DDT

- Linux:
 - Linux IA32 (Intel and AMD)
 - Linux IA64
 - Linux Opteron
- Unix
 - PowerPC (AIX)
 - SGI Altix
 - SGI Irix
 - SUN Sparc
 - PA-Risc and Itanium Superdome

Availability of DDT at HWW

Platform	Availability	Remarks
Volvox	Yes	V 1.4
AzusA	Yes	V 1.4

More information:

<http://www.hlrs.de/organization/tsc/services/tools/debugger/ddt/>

Availability of DDT at University Stuttgart

- Two user Floating License for University Stuttgart:
 1. Download Software from <http://www.etnus.com>
 2. Set environment variable for license.

`LM_LICENSE_FILE=7244@servint1.rus.uni-stuttgart.de`

More information about campus licenses available at

<http://www.hlrs.de/organization/par/services/tools/campus>

DDT usage at HLRS

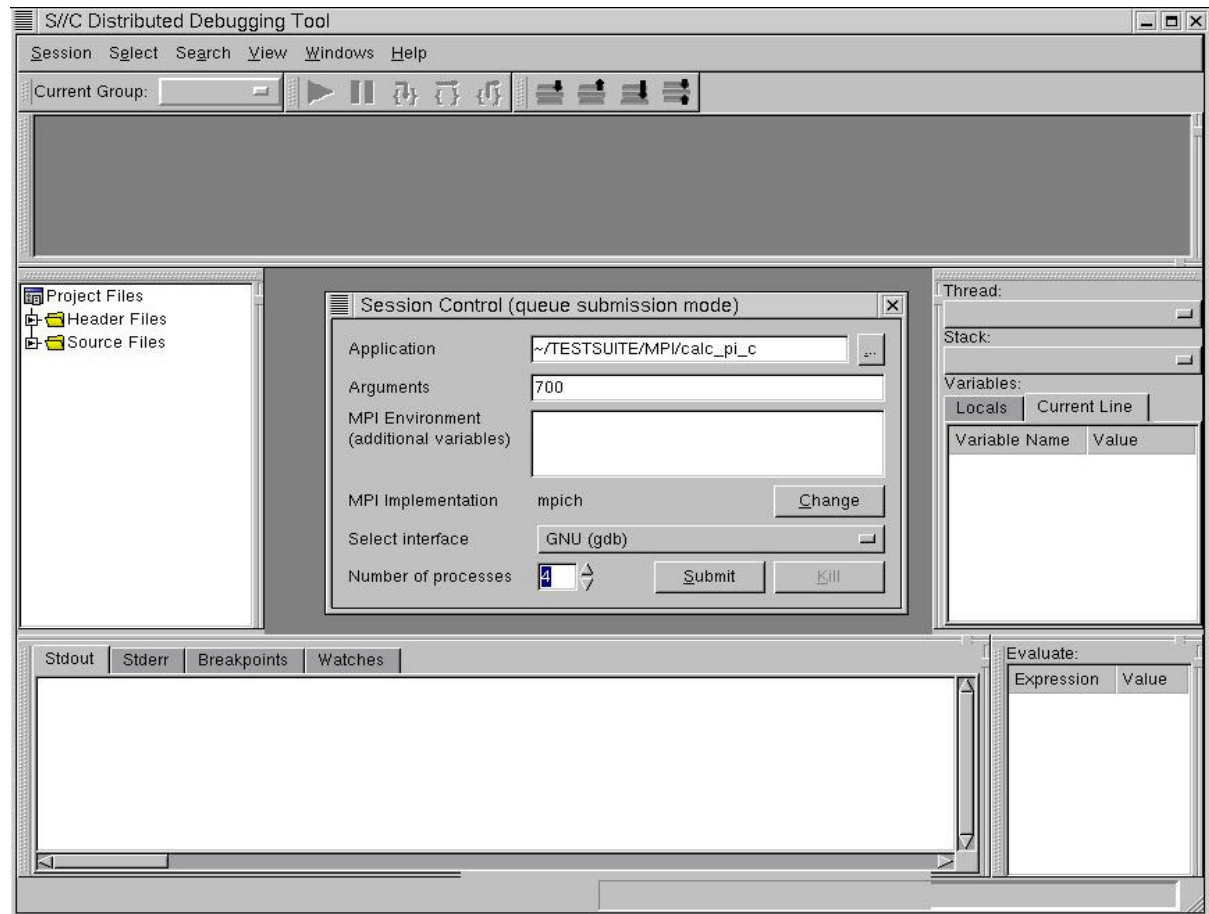
- Set USE_DDT in your login scripts
- Compile with -g compiler switch
- Command name: **ddt** or **\$DDT**
- To start debugging with DDT simply type:
% **\$DDT myprog arguments to myprog**

DDT Look & Feel

DDT Main
Window

Configuration
Window

and all
belonging
Panels
(Thread,
Stack,
Output,
Source code,
etc.)



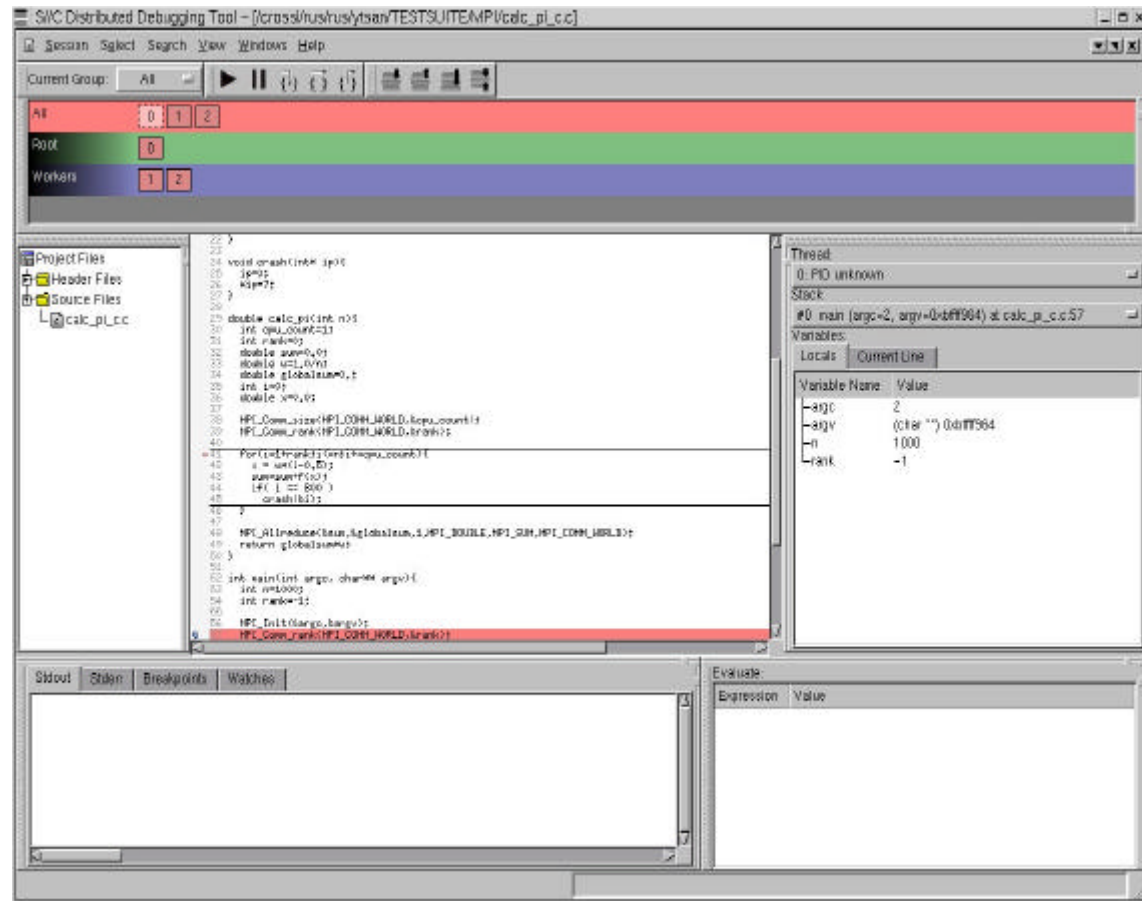
DDT Main/Process Window

MPI
Groups

File
browse
and
Source
pane

Output,
Breakpo
ints,
Watch

Pane

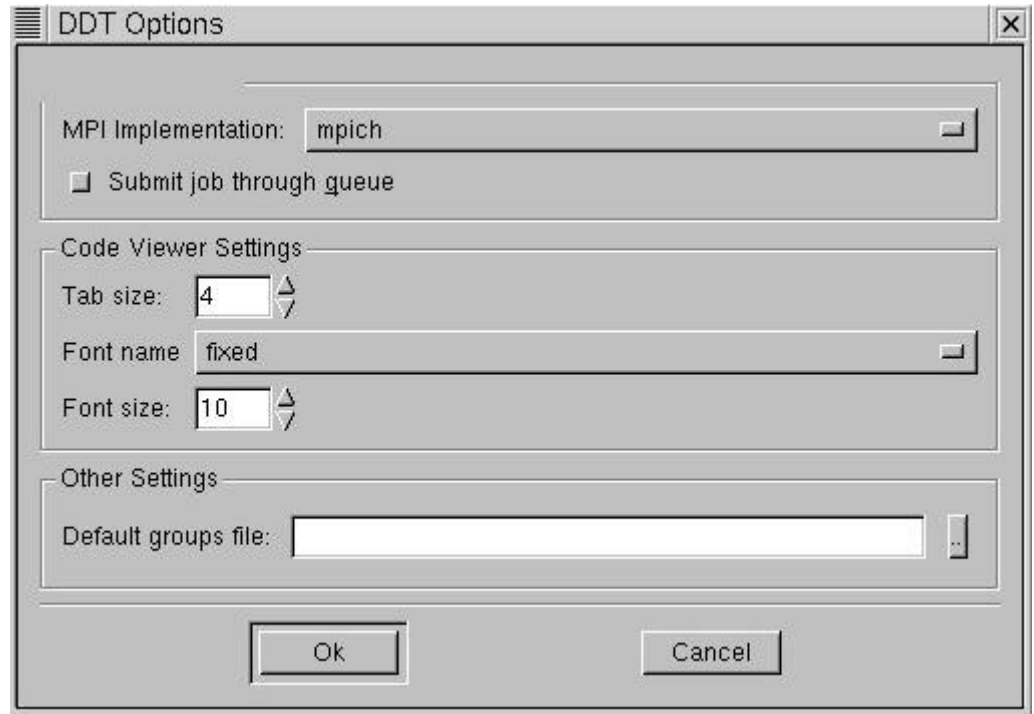


Thread,
Stack,
Local
and
Global
Variables
Pane
Evaluation
window

DDT Options Window

The Options window:
here is the MPI
library implementation
selected

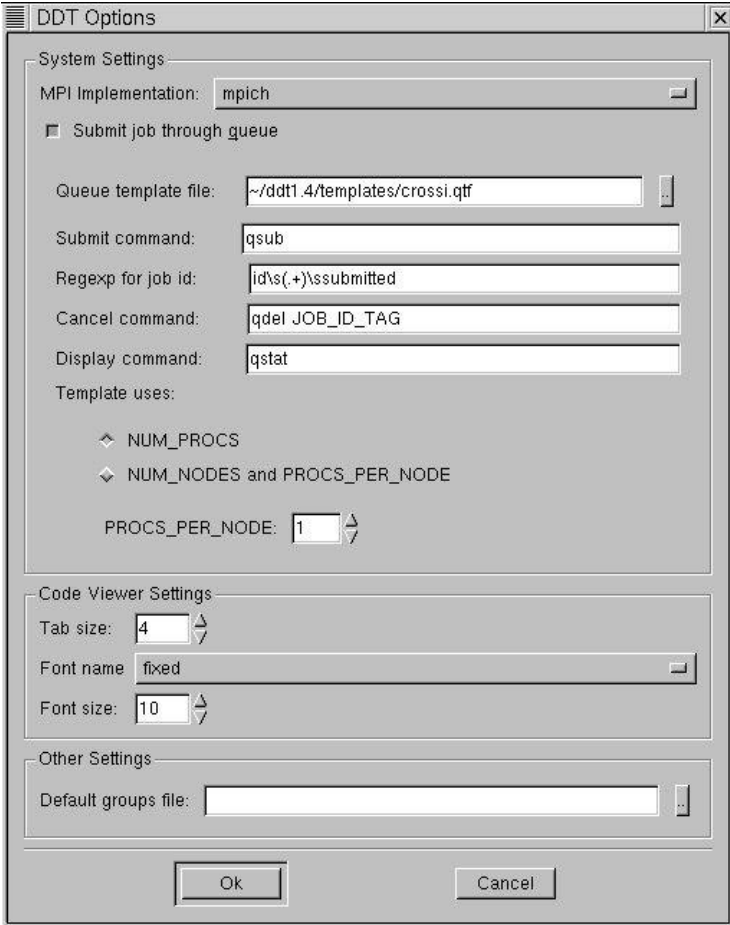
The program can also
be submitted through
some batch system



DDT Options Window (Queue)

The Options windows:
This program uses mpich
as MPI implementation and
starts the program through
PBS batch system

For more information
about the Template files
and the commands,
please read the DDT
Users Manual.



The screenshot shows the 'DDT Options' dialog box with the 'Queue' tab selected. The 'System Settings' section includes a dropdown for 'MPI Implementation' set to 'mpich', a checkbox for 'Submit job through queue' which is checked, and several text fields for 'Queue template file' (~/ddt1.4/templates/crossi.qtf), 'Submit command' (qsub), 'Regexp for job id' (id\s(.+)\ssubmitted), 'Cancel command' (qdel JOB_ID_TAG), and 'Display command' (qstat). Below these is a 'Template uses:' section with expandable items for 'NUM_PROCS' and 'NUM_NODES and PROCS_PER_NODE', and a 'PROCS_PER_NODE' spinner set to 1. The 'Code Viewer Settings' section has 'Tab size' (4), 'Font name' (fixed), and 'Font size' (10). The 'Other Settings' section has a 'Default groups file' field. 'Ok' and 'Cancel' buttons are at the bottom.

Section	Parameter	Value
System Settings	MPI Implementation	mpich
	Submit job through queue	<input checked="" type="checkbox"/>
	Queue template file	~/ddt1.4/templates/crossi.qtf
	Submit command	qsub
	Regexp for job id	id\s(.+)\ssubmitted
	Cancel command	qdel JOB_ID_TAG
Code Viewer Settings	Tab size	4
	Font name	fixed
Code Viewer Settings	Font size	10
	Other Settings	Default groups file



DDT Thread and Stack Window

Thread Pane:

Switch between all program threads

Stack Pane:

Switch between the functions in the selected thread

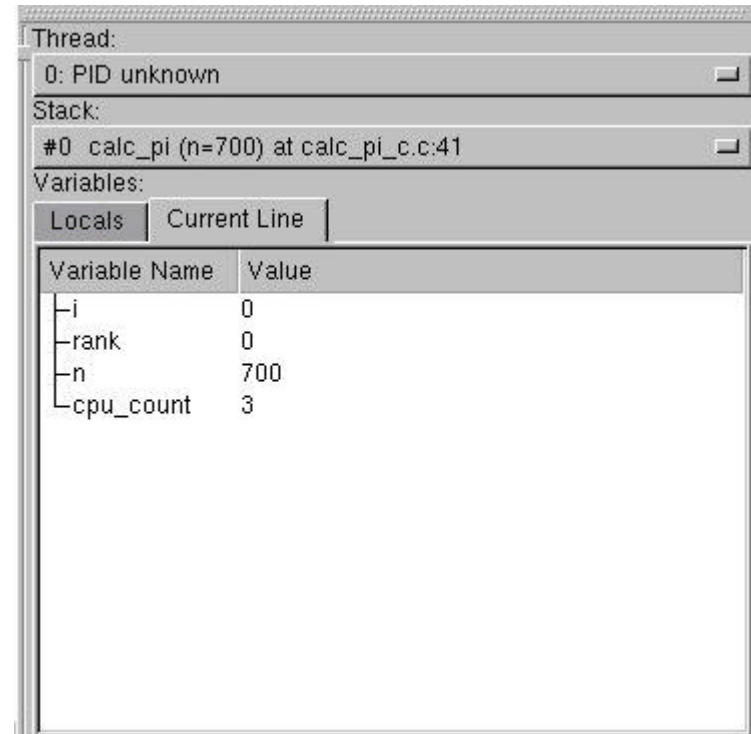
Variables Pane:

Local variables:

Shows the variables value for the function in which ddt is currently stopped

Current Line:

Shows every variable value between the two lines selected by the user



DDT Source Pane

With the right mouse button Set or Remove a breakpoint at the selected line

When the program is running and stopped at a breakpoint, the line is coloured in red (by OpenMP programs) and red, blue or green by programs using MPI

```
3 // Simple OpenMP Program to calculate PI
4 //
5 //
6 // Author: Matthias Mueller <mueller@hirs.de>
7 // Thu Dec 28 15:12:09 CET 2000
8 //
9 // Usage: calc_pi [iterations]
10 // the program will crash if iterations >= 800
11 //
12 // This is the intended behavior. Because the program
13 // test a debugger.
14 //
15
16 #include <stdio.h>
17 #include <stdlib.h>
18
19 double f(double a){
20     return 4.0/(1.0 + a*a);
21 }
22
23 void crash(int* ip){
24     ip=0;
25     *ip=7;
26 }
27
28 double calc_pi(int n){
29     double sum=0.0;
30     double w=1.0/n;
31     int i=0;
32     double x=0.0;
33     #pragma omp parallel for reduction(+:sum)
34     for(i=1;i<=n;i++){
35         x = w*(i-0.5);
36         sum=sum+f(x);
37         if( i == 800 )
38             crash(&i);
39     }
40     return sum*w;
41 }
42
43 int main(int argc, char** argv){
44     int n=1000;
45     if(argc>2){
46         fprintf(stderr," Usage : %s [iterations] \n",argv[0]);
47         return 1;
48     }
49     if(argc==2){
50         n=atoi(argv[1]);
51     }
52     printf(" Pi = %g\n",calc_pi(n));
53     return 0;
54 }
```

Parallel Debugging - Philosophy

- By default, DDT places processes in groups
 - All Group - Includes parent and all related processes
 - Root/Workers Group - Only processes that share the same source code
- Command can act on single process or group
 - stop process , stop group
 - next step process , next step group
 - go process, go group

DDT more information

- http://www.streamline-computing.com/softwaredivision_1.shtml