





# Debugging with Totalview

*Martin Čuma*

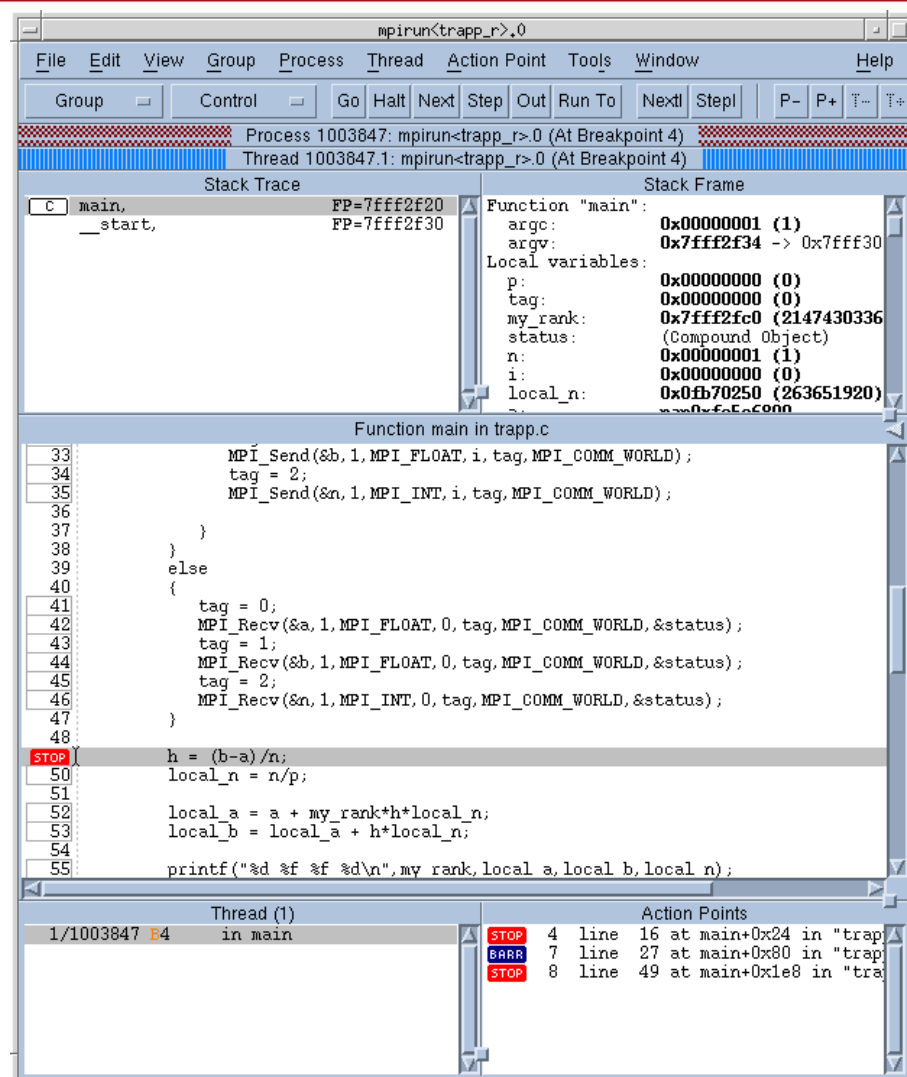
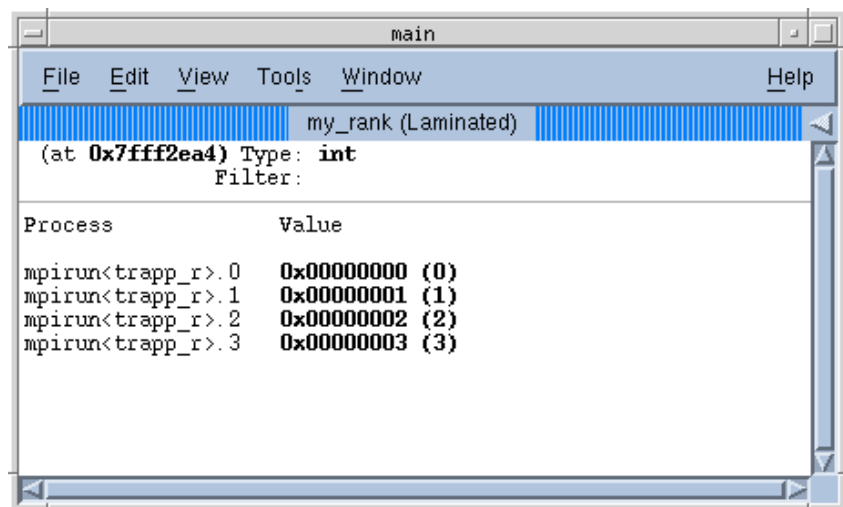
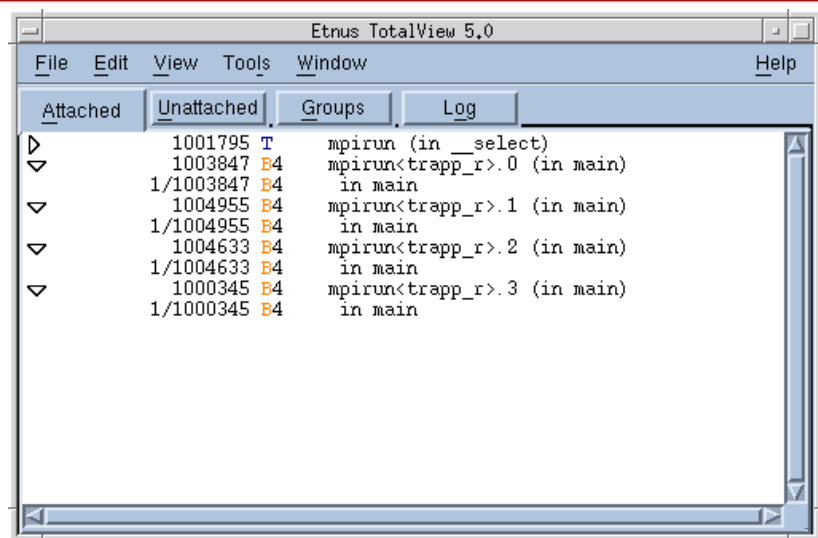
*Center for High Performance  
Computing University of Utah  
mcuma@chpc.utah.edu*




- Totalview introduction.
- Basic operation.
- Serial debugging.
- Parallel debugging.
- OpenMP
- MPI
- Memory debugging.
- Some other useful information




- source and machine level debugger
- command line and graphic interface
- serial and parallel debugging support
- supports remote debugging
- runs on variety of platforms





- Data examination
  - view data in the variable windows
  - change the values of variables
  - modify display of the variables
  - visualize data
- Action points
  - breakpoints and barriers (static or conditional)
  - watchpoints
  - evaluation of expressions



- Automatic attachment of child processes
- Create process groups
- Share breakpoints among processes
- Process barrier breakpoints
- Process group single-stepping
- “Laminate” variables
- Display MPI message queue state



## 1. Compile binary with debugging information

- flag -g

```
g77 -g test.f -o test
```

- if use fork() or execve(), link ...

```
g77 -g -L/.../totalview/linux-x86-64/lib -ldbfork  
test.f -o test
```


## 2. Run Totalview

- TV + executable

```
totalview executable
```

- TV + core file

```
totalview executable core_file
```



- run TV and attach the executable
  - start TV
  - menu New Program Window
  - fill in executable file name
- run TV and attach running program
  - start TV
  - menu New Program Window
  - pullout Attach to an Existing Process
  - choose process ID and fill in executable file name

### 3. Totalview operation

- left mouse button - select
- right mouse button - menu
- left mouse button double click - dive






- CUSP – DFT plane wave atomic simulation
  - atomic nuclei interact classically
  - electrons described by wavefunctions - plane waves
  - interaction between w.f.'s and nuclei define the system potential and kinetic energy
  - iterate electronic and nuclear energy to self consistency to achieve ground state
- Implementation
  - 10+ Fortran source code files
  - simple Makefile w/ flags **-g -byteswapio**

- Stack trace – procedure hierarchy
- Stack frame – variables display
- Source code – code + process navigation
- Threads list – in case of multithreaded application
- Action points – list of breakpoints, barriers,...

- Menu Go/Halt/Next/Step/Hold or shortcuts
- Possible actions (thread,process/group):
  - go (g/G)
  - halt (h/H)
  - step (source line) (s/S)
  - step (instruction) (i/I)
  - next (source line) (n/N)
  - next (instruction) (x/X)
  - run (to selection) (r/R)
  - return (out of function) (o/O)

- Breakpoints and barriers
  - toggle location with left mouse (shift for barrier)
  - right-click – Properties for options
- Evaluation points
  - set conditional breakpoints
  - conditionally patch out code
- Watchpoints
  - watch for change in a memory location

- Variable view
  - dive (right mouse) on any variable
  - change data type
  - select an array slice, e.g. (3:3,:)
  - filter array values, e.g. .ne. 0
- Variable visualization
- menu Visualize – only up to 2D arrays



- Compilation

- Arches:

```
pathf90 -openmp test.f -g -o test.exe
```

```
pgf90 -mp test.f -g -o test.exe
```

```
ifort -openmp test.f -g -o test.exe
```


```
gcc4 -fopenmp test.c -g -o test.exe
```

- Running


- set OMP\_NUM\_THREADS

- Example – saxpy routine

- simple vector-scalar product




- TV automatically attaches all threads
- put breakpoint to OpenMP parallel section to debug threads
- variable lamination - show values from all threads in one window – does not always work
- barrier points – shift-left click
- ambiguous action points – select all




- Gravity and magnetic field inversion
  - Just a code that I am working on now
- Implementation
  - 1D domain decomposition
  - Each process has some slices of the 3D domain
  - MPI used for data synchronization, etc



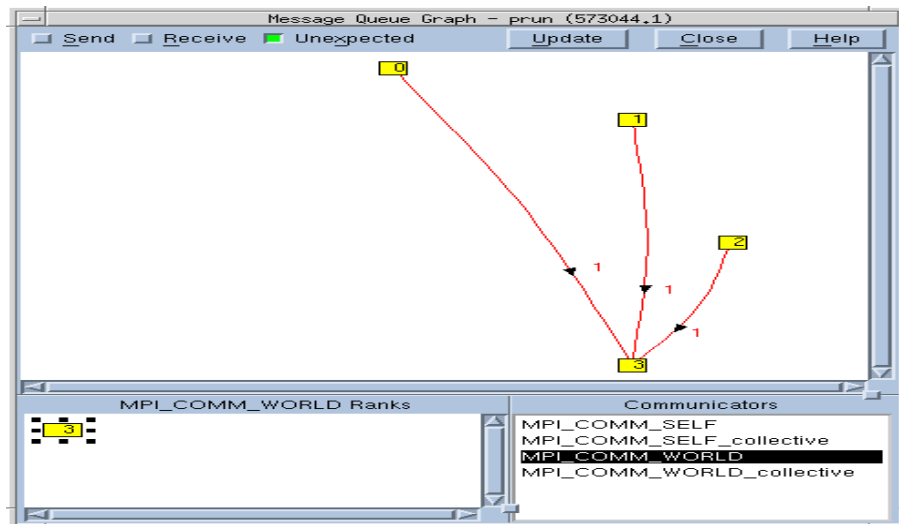


- MPICH2 on interactive nodes
  - good enough for most parallel problems
  - **source**  
`/uufs/chpc.utah.edu/sys/pkg/mpich2/std_pgi/etc/mpich2.csh`
  - `mpicc submit.o -g -o submit`
- Running
  - Start MPD daemon on the interactive node  
`mpdboot -n 1`  
if have problems, create `~/.mpd.conf` file
  - Run:  
open Totalview (**totalview**), then in New Program Window enter:
    - program name
    - MPICH2 for parallel sysem
    - number of processors




- Running on compute nodes
  - use larger number of processors or other MPI than MPICH2 is needed
  - `source ../etc/mpi*.csh`
  - `mpicc -g -mp submit.c`
  - `mpicc submit.o -o submit`
- Running on compute nodes (large # of processors)
  - CHPC has 32 process Totalview license
  - interactive PBS:  
`qsub -I -X -l nodes=4:ppn=4,walltime=2:00:00`


- Process synchronization – program groups
- Barrier points
- Message queue state graph and display




submit_ib.0	
File Edit View Window Help	
Message State for "submit_ib.0" (20534.1)	
MPI_COMM_WORLD	
Comm_size	2
Comm_rank	0
Pending receives	: none
Unexpected messages	: none
Pending sends	: none
MPI_COMM_WORLD_collective	
Comm_size	2
Comm_rank	0
Pending receives	: none
Unexpected messages	: none
Pending sends	: none
MPI_COMM_SELF	
Comm_size	1
Comm_rank	0
Pending receives	: none
Unexpected messages	: none
Pending sends	: none
MPI_COMM_SELF_collective	
Comm_size	1
Comm_rank	0
Pending receives	: none
Unexpected messages	: none
Pending sends	: none




- Dynamic memory debugging tool
- display memory status
- paint allocated and deallocated blocks
- find memory leaks
- identify dangling pointers
- enable with **Tools > Memory Debugger  
> Enable memory debugging checkbox**



- Allows to reversely debug the code
- Must be turned on at the start of debugging session
- Run to the error, then backtrack to the source of the problem
- Helps to capture race conditions and other hard to reproduce bugs



- Debug Nvidia CUDA on GPU
- We haven't tested it but according to some users it works



- TotalviewTech webpage

`http://www.roguewave.com/products/totalview`

- Location of Totalview

`Arches:/uufs/chpc.utah.edu/sys/pkg/totalview/std`

Some group desktops: `inquire` at CHPC

- Documentation

`http://www.roguewave.com/support/product-documentation/totalview.aspx`

`http://www.chpc.utah.edu/software/docs/par\_devel.html`

`http://www.chpc.utah.edu/software/docs/totalview.html`

`http://www.chpc.utah.edu/short\_courses/Totalview`


- Free for students
- Limited to one computer, 4 processes
- To sign up, e-mail [mcuma@chpc.utah.edu](mailto:mcuma@chpc.utah.edu):
  - name
  - e-mail
  - university ID
  - anticipated year of graduation
- More details

[http://www.totalviewtech.com/academia/totalview\\_express.html](http://www.totalviewtech.com/academia/totalview_express.html)




- compilers check for syntax errors
  - some compiler flags help too (-C)
- memory checking tools - many errors are due to bad memory management
  - valgrind – easy to use
  - purify – harder to use

- Good for finding OpenMP errors
  - race conditions
  - privatization
- Intel thread checker (ITC)
  - OpenMP
  - pthreads



- We have a free classroom license
- Tools for all stages of development
  - Compilers and libraries
  - Verification tools
  - Profilers
- More info

<http://software.intel.com/en-us/intel-sdp-home/>



- Thread checking
  - Data races and deadlocks
- Memory checker
  - Like leaks or corruption
- Standalone or GUI integration
- More info

<http://software.intel.com/en-us/articles/intel-inspector-xe/>

- **Source the environment**

```
source /uufs/chpc.utah.edu/sys/pkg/intel/  
composerxe-2011.1.107/bin/compilervars.csh intel64  
addpath "/uufs/chpc.utah.edu/sys/pkg/intel/inspector_xe_2011/  
bin64:$PATH"
```

- **Compile with -tcheck -g**

```
ifort -openmp -tcheck -g trap.f
```

- **Run tcheck**

`inspxe-gui` – graphical user interface

`inspxe-cl` – command line

- **Tutorial**

<http://software.intel.com/sites/products/documentation/hpc/inspectorxe/en-us/lin/start/index.htm>

- Serial and thread profiler
- Built in hardware counters
  - Find hardware related slowdowns

- Source the environment

```
addpath "/uufs/chpc.utah.edu/sys/pkg/intel/  
vtune_amplifier_xe_2011/bin64:$PATH"
```

- Run Vtune

`amplxe-gui` – graphical user interface

`amplxe-cl` – command line

- More info

<http://software.intel.com/en-us/articles/intel-vtune-amplifier-xe/>



- We have license for more-less everything
- Installed in the /uufs/chpc.utah.edu branch
- Visible on all clusters and Linux desktops
- See the products description and documentation

<http://software.intel.com/en-us/articles/intel-vtune-amplifier-xe/>

- Let us know if you are interested in hands on tutorial